

BEN-GURION UNIVERSITY OF THE NEGEV

DATA STRUCTURES

202.1.1031

Assignment No. 4 - Solution

Author:

Publish date: 20.5.2023

Submission date: 10.06.2023

אוניברסיטת בן-גוריון בנגב
Ben-Gurion University of the Negev



Contents

| | | |
|----------|---|-----------|
| 2 | Skip-List | 3 |
| 2.1 | Warm-Up and Familiarization | 3 |
| 2.1.1 | Implementation of Abstract Functions | 3 |
| 2.1.2 | Analysis of the Probabilistic Process | 3 |
| 2.1.3 | Analysis of the operations | 4 |
| 2.2 | Order Statistics | 7 |
| 3 | Hashing | 8 |
| 3.1 | Introduction | 8 |
| 3.1.1 | Hash Functions | 8 |
| 3.2 | Hash Implementations | 8 |
| 3.3 | Hash Tables | 8 |
| 3.3.1 | Introduction | 8 |
| 3.4 | Theoretical Questions | 9 |
| 4 | Designing a data structure according to given specifications | 11 |

Tasks and Questions

| | | | |
|------|--------|-----|----|
| 2.1 | Answer | ... | 3 |
| 2.2 | Answer | ... | 3 |
| 2.3 | Answer | ... | 3 |
| 2.4 | Answer | ... | 4 |
| 2.5 | Answer | ... | 4 |
| 2.6 | Answer | ... | 4 |
| 2.7 | Answer | ... | 5 |
| 2.8 | Answer | ... | 6 |
| 2.9 | Answer | ... | 6 |
| 2.10 | Answer | ... | 6 |
| 2.11 | Answer | ... | 7 |
| 2.12 | Answer | ... | 7 |
| 3.1 | Answer | ... | 8 |
| 3.2 | Answer | ... | 8 |
| 3.3 | Answer | ... | 8 |
| 3.4 | Answer | ... | 8 |
| 3.5 | Answer | ... | 8 |
| 3.6 | Answer | ... | 8 |
| 3.7 | Answer | ... | 8 |
| 3.8 | Answer | ... | 8 |
| 3.9 | Answer | ... | 8 |
| 3.10 | Answer | ... | 8 |
| 3.11 | Answer | ... | 9 |
| 3.12 | Answer | ... | 9 |
| 3.13 | Answer | ... | 9 |
| 3.14 | Answer | ... | 9 |
| 3.15 | Answer | ... | 10 |
| 3.16 | Answer | ... | 10 |
| 4.1 | Answer | ... | 11 |

2 Skip-List

2.1 Warm-Up and Familiarization

2.1.1 Implementation of Abstract Functions

Answer 2.1: [Implementation in code](#)

2.1.2 Analysis of the Probabilistic Process

Answer 2.2: The tables are:

| p = 0.33 | | | | | | | |
|----------|----------------|----------------|----------------|----------------|----------------|------------------------------|---|
| x | $\hat{\ell}_1$ | $\hat{\ell}_2$ | $\hat{\ell}_3$ | $\hat{\ell}_4$ | $\hat{\ell}_5$ | Expected Level ($E[\ell]$) | Average delta ($\frac{1}{5} \cdot \sum_{i=1}^5 (\hat{\ell}_i - E[\ell])$) |
| 10 | 3.000 | 4.200 | 2.800 | 3.100 | 4.1000 | 3.030 | 0.514 |
| 100 | 2.870 | 2.860 | 3.150 | 3.250 | 3.420 | 3.030 | 0.212 |
| 1000 | 3.057 | 2.997 | 3.173 | 3.017 | 3.167 | 3.030 | 0.071 |
| 10000 | 2.991 | 3.011 | 3.019 | 3.045 | 3.015 | 3.030 | 0.020 |

| p = 0.5 | | | | | | | |
|---------|----------------|----------------|----------------|----------------|----------------|------------------------------|---|
| x | $\hat{\ell}_1$ | $\hat{\ell}_2$ | $\hat{\ell}_3$ | $\hat{\ell}_4$ | $\hat{\ell}_5$ | Expected Level ($E[\ell]$) | Average delta ($\frac{1}{5} \cdot \sum_{i=1}^5 (\hat{\ell}_i - E[\ell])$) |
| 10 | 2.400 | 1.700 | 2.300 | 2.200 | 2.000 | 2.000 | 0.240 |
| 100 | 2.000 | 2.080 | 1.840 | 2.040 | 2.060 | 2.000 | 0.068 |
| 1000 | 2.009 | 2.067 | 2.048 | 1.962 | 2.065 | 2.000 | 0.045 |
| 10000 | 1.987 | 2.016 | 2.011 | 2.012 | 1.993 | 2.000 | 0.012 |

| p = 0.75 | | | | | | | |
|----------|----------------|----------------|----------------|----------------|----------------|------------------------------|---|
| x | $\hat{\ell}_1$ | $\hat{\ell}_2$ | $\hat{\ell}_3$ | $\hat{\ell}_4$ | $\hat{\ell}_5$ | Expected Level ($E[\ell]$) | Average delta ($\frac{1}{5} \cdot \sum_{i=1}^5 (\hat{\ell}_i - E[\ell])$) |
| 10 | 1.600 | 1.300 | 1.800 | 1.700 | 1.300 | 1.333 | 0.233 |
| 100 | 1.290 | 1.270 | 1.290 | 1.530 | 1.280 | 1.333 | 0.080 |
| 1000 | 1.347 | 1.358 | 1.306 | 1.309 | 1.337 | 1.333 | 0.019 |
| 10000 | 1.333 | 1.332 | 1.324 | 1.331 | 1.330 | 1.333 | 0.003 |

| p = 0.9 | | | | | | | |
|---------|----------------|----------------|----------------|----------------|----------------|------------------------------|---|
| x | $\hat{\ell}_1$ | $\hat{\ell}_2$ | $\hat{\ell}_3$ | $\hat{\ell}_4$ | $\hat{\ell}_5$ | Expected Level ($E[\ell]$) | Average delta ($\frac{1}{5} \cdot \sum_{i=1}^5 (\hat{\ell}_i - E[\ell])$) |
| 10 | 1.100 | 1.200 | 1.500 | 1.000 | 1.000 | 1.111 | 0.142 |
| 100 | 1.150 | 1.100 | 1.130 | 1.220 | 1.140 | 1.111 | 0.041 |
| 1000 | 1.118 | 1.115 | 1.108 | 1.110 | 1.100 | 1.111 | 0.005 |
| 10000 | 1.116 | 1.109 | 1.109 | 1.113 | 1.110 | 1.111 | 0.003 |

Answer 2.3: The higher the probability for success the lower the average number of levels. (P is in inverse proportion to the average height)

Answer 2.4: The larger the sample size (x) we are using the lower the delta we get. This is because over a lot of operations the random elements even out.

2.1.3 Analysis of the operations

Answer 2.5: [Implementation in code](#)

Answer 2.6: The tables are:

| $p = 0.33$ | | | |
|------------|-------------------|----------------|------------------|
| x | Average Insertion | Average Search | Average Deletion |
| 1000 | 425.393 | 157.68 | 71.093 |
| 2500 | 368.808 | 178.936 | 36.388 |
| 5000 | 305.593 | 195.469 | 35.445 |
| 10000 | 331.812 | 214.872 | 40.699 |
| 15000 | 378.559 | 247.66 | 43.532 |
| 20000 | 359.962 | 269.007 | 50.1 |
| 50000 | 437.585 | 402.648 | 75.401 |

| $p = 0.5$ | | | |
|-----------|-------------------|----------------|------------------|
| x | Average Insertion | Average Search | Average Deletion |
| 1000 | 263.62 | 127.465 | 20.273 |
| 2500 | 224.68 | 160.755 | 20.38 |
| 5000 | 254.393 | 191.164 | 23.582 |
| 10000 | 290.292 | 231.926 | 36.012 |
| 15000 | 317.342 | 260.408 | 36.786 |
| 20000 | 321.419 | 269.433 | 37.494 |
| 50000 | 403.929 | 428.393 | 62.658 |

| $p = 0.75$ | | | |
|------------|-------------------|----------------|------------------|
| x | Average Insertion | Average Search | Average Deletion |
| 1000 | 172.38 | 128.478 | 17.753 |
| 2500 | 201.575 | 172.668 | 15.352 |
| 5000 | 232.879 | 215.858 | 17.533 |
| 10000 | 280.273 | 258.018 | 24.16 |
| 15000 | 299.252 | 297.724 | 26.178 |
| 20000 | 322.698 | 313.567 | 27.05 |
| 50000 | 420.842 | 498.215 | 50.269 |

| $p = 0.9$ | | | |
|-----------|-------------------|----------------|------------------|
| x | Average Insertion | Average Search | Average Deletion |
| 1000 | 210.327 | 156.66 | 11.507 |
| 2500 | 255.7 | 229.199 | 13.279 |
| 5000 | 331.055 | 303.192 | 21.68 |
| 10000 | 369.724 | 365.948 | 21.682 |
| 15000 | 416.744 | 427.251 | 27.954 |
| 20000 | 449.097 | 501.986 | 27.703 |
| 50000 | 686.795 | 908.683 | 67.923 |

Answer 2.7: The graph are:





Answer 2.8: Insertion The average insertion time is the lowest when P is around the halfway point between 0 and 1. that because when P is too high the generate-height function and thus the linking on each level takes more time (because there are more levels). And also the find function used in the insert operation take longer due to the list being higher (that's because we also need to go down as well as to the right).

And if the P value is too low there won't be enough variety in different nodes' heights, making the find function more and more similar to linear search - thus making it slower.

Search The lower the P value, the lower the average search time, up to a certain threshold, due to the same reason as in Insertion.

Delete The lower the P, the higher the average deletion time.

Answer 2.9: Let SL be a skip-list with n elements.

The probability that an element x is in the list S_i is P^i , And the size of S_i has binomial distribution with parameters n and P^i . So $E[|S_i|] = n * P^i$

Therefore, the expected number of nodes is: $E[\sum_{i=0}^{\infty} |S_i|]$ and using the linearity of expectation we get $\sum_{i=0}^{\infty} E[|S_i|] = \sum_{i=0}^{\infty} n * P^i = n * \sum_{i=0}^{\infty} P^i$ and using the formula for sum of infinite geometric series with $|quotient| < 1$ we get $n * \frac{1}{1-P} = \frac{n}{1-P}$

Thus the expected number of nodes in a conceptual Skip-List containing elements, not including the Sentinels is $\frac{n}{1-P}$

Answer 2.10: For each $i > 0$ define the indicator random variable:

$$X_i = \begin{cases} 0, & \text{if } S_i \text{ is empty} \\ 1, & \text{if } S_i \text{ is not empty} \end{cases} \quad (1)$$

The height h of the skip list is then given by: $h = \sum_{i=0}^{\infty} X_i$

Note that $X_i \leq |S_i|$, thus $E[X_i] \leq E[|S_i|] = n * P^i$

Additionally, $X_i \leq 1$ so $E[X_i] \leq 1$

$$E[h] = E[\sum_{i=1}^{\infty} X_i] = \sum_{i=1}^{\infty} E[X_i]$$

$$\text{Auxiliary calculation: } n * P^i = 1 \Rightarrow P^i = \frac{1}{n} \Rightarrow i = \log_P(\frac{1}{n})$$

$$\begin{aligned} \sum_{i=1}^{\infty} E[X_i] &= \sum_{i=1}^{\lfloor \log_P(\frac{1}{n}) \rfloor} E[X_i] + \sum_{\lfloor \log_P(\frac{1}{n}) \rfloor + 1}^{\infty} E[X_i] \\ &\leq \sum_{i=1}^{\lfloor \log_P(\frac{1}{n}) \rfloor} 1 + \sum_{\lfloor \log_P(\frac{1}{n}) \rfloor + 1}^{\infty} n * P^i \end{aligned}$$

Auxiliary calculation: $\sum_{i=0}^{\lfloor \log_P(\frac{1}{n}) \rfloor + 1} n * P^i = n * P^{\lfloor \log_P(\frac{1}{n}) \rfloor + 1} + n * P^{\lfloor \log_P(\frac{1}{n}) \rfloor + 2} + \dots =$
 $n * P^{\lfloor \log_P(\frac{1}{n}) \rfloor + 1} * [\sum_{i=0}^{\infty} P^i] = n * P^{\lfloor \log_P(\frac{1}{n}) \rfloor + 1} * \frac{1}{1-P} \leq n * P^{\log_P(\frac{1}{n})} * \frac{1}{1-P} = n * \frac{1}{n} * \frac{1}{1-P} = \frac{1}{1-P}$

$$\sum_{i=1}^{\lfloor \log_P(\frac{1}{n}) \rfloor} 1 + \sum_{i=\lfloor \log_P(\frac{1}{n}) \rfloor + 1}^{\infty} n * P^i = \log_P(\frac{1}{n}) + \frac{1}{1-P}$$

So the expected height of a Skip-List as a function of p is most $\log_P(\frac{1}{n}) + \frac{1}{1-P}$

Answer 2.11: As shown in lectures, $E[X] = \frac{1}{P}$.

Let R_i denote the number of steps the path has at S_i .

The number of steps of a reverse path is some S_i is a random variable which is similar to a geometric distribution random variable. except that we don't count the last coin flip and we stop when reaching a sentinel, so $R_i \leq X_i - 1$.

Using linearity of expectation we get: $E[R_i] \leq E[X_i - 1] = E[X_i] - 1 = \frac{1}{1-P} - 1 = \frac{P}{P-1}$.

Additionally, $R_i \leq |S_i|$ by the way it was defined, therefore: $E[R_i] \leq E[|S_i|] = n * P^i$.

Let R by the length of the search path for some search value.
 $R = h + \sum_{i=0}^{\infty} R_i$.

Auxiliary calculation: $n * P^i = 1 \Rightarrow P^i = \frac{1}{n} \Rightarrow i = \log_P(\frac{1}{n})$

$$\begin{aligned} \text{So we'll get } E[R] &= E[h + \sum_{i=0}^{\infty} R_i] = E[h] + \sum_{i=0}^{\infty} E[R_i] = E[h] + \sum_{i=0}^{\lfloor \log_P(\frac{1}{n}) \rfloor} E[R_i] + \sum_{i=\lfloor \log_P(\frac{1}{n}) \rfloor + 1}^{\infty} X_i \\ &\leq E[h] + \sum_{i=0}^{\lfloor \log_P(\frac{1}{n}) \rfloor} \frac{P}{1-P} + \sum_{i=\lfloor \log_P(\frac{1}{n}) \rfloor + 1}^{\infty} n * P^i \\ &\leq \log_P(\frac{1}{n}) + \frac{1}{1-P} + (\log_P(\frac{1}{n}) + 1) * (\frac{P}{1-P}) + \frac{1}{1-P} \\ &= \log_P(\frac{1}{n}) + \frac{1}{1-P} + \frac{(P * \log(\frac{1}{n})) + P}{1-P} + \frac{1}{1-P} \\ &= \frac{\log_P(\frac{1}{n}) + 2 + P}{1-P} \end{aligned}$$

Thus the expected time complexity of Insertion/Search as a function of p is: $\frac{\log_P(\frac{1}{n}) + 2 + P}{1-P}$

Comparing this to the result we've got in task 2.6 it seems that the result are obeying this formula; as very high P value (close to 1) bump the time due to the fact that the numerator increase and the denominator decreases 1-P.

2.2 Order Statistics

Answer 2.12: Sorry I didn't have enough time. this one is a bit unfair.

3 Hashing

3.1 Introduction

3.1.1 Hash Functions

Answer 3.1: After we perform true Modulus operation the result is always non-negative. and thus the sign bit of the result is 0.

\ggg is unsigned-shift, so performing it $(w - k)$ times result in shifting everything to the right $(w - k)$ times and filling it with 0s.

And when looking at a binary number, shifting it to the right is equivalent to dividing it by 2 (integer division) $(w - k)$ times, or $2^{(w-k)}$ - thus the two equation are equals.

Sources: <https://stackoverflow.com/questions/2811319/difference-between-and>

<https://stackoverflow.com/questions/5385024/mod-in-java-produces-negative-numbers>

Answer 3.2: We can look at the object's binary representation and collect it into chunks of 8 bits (fill the reminder of the last one with 0s if needed). After that look at each chunk as a single ASCII character and use the Carter-Wegman hashing for strings on the resulted string.

3.2 Hash Implementations

Answer 3.3: [Implementation in code](#)

Answer 3.4: [Implementation in code](#)

Answer 3.5: [Implementation in code](#)

3.3 Hash Tables

3.3.1 Introduction

Answer 3.6: [Implementation in code](#)

Answer 3.7: [Implementation in code](#)

Answer 3.8: The results are:

| Linear Probing | | |
|----------------|-------------------|----------------|
| $\max \alpha$ | Average Insertion | Average Search |
| 1/2 | 38.662 | 40.893 |
| 3/4 | 33.784 | 57.281 |
| 7/8 | 34.242 | 95.738 |
| 15/16 | 40.789 | 287.951 |

Answer 3.9: The higher the load factor, the slower each operation takes, and it approach $O(n)$ as the load factor approach 1. That's because in each traversal of the table we need to look for an empty cell for longer time (because the table is in higher capacity).

Answer 3.10: The results are:

| Chaining | | |
|--------------|-------------------|----------------|
| max α | Average Insertion | Average Search |
| 1/2 | 101.138 | 47.909 |
| 3/4 | 55.848 | 36.738 |
| 1 | 55.413 | 42.524 |
| 3/2 | 62.837 | 54.872 |
| 2 | 61.570 | 57.587 |

Answer 3.11: The higher the load factor the higher the search on average, that's because in addition to going into the right cell we also need to perform linear search in the linked list to find the item (if exist). But the insertion stays relatively constant because we insert element to the head of the list - making its length irrelevant.

Answer 3.12: The results are:

| Long operations | | |
|-----------------|-------------------|----------------|
| max α | Average Insertion | Average Search |
| 1 | 91.441 | 40.1893 |

| String operations | | |
|-------------------|-------------------|----------------|
| max α | Average Insertion | Average Search |
| 1 | 9922.089 | 9993.572 |

We can deduce that hashing strings is a lot more expensive than hashing numbers using the Dietzfelbinger et al hashing algorithm.

3.4 Theoretical Questions

Answer 3.13: The algorithm for finding the Successor is a brute force over all the elements of the hash table (one list at a time) until we find smallest number larger than val (aka the successor). Because we go over the entire data-structure one item at a time the complexity is $O(n)$

Require: val

```

1: current - value  $\leftarrow \infty$ 
2: for map[0] to map[map.size() - 1] do
3:   list  $\leftarrow$  map[i]
4:   for element in list do
5:     if (element < current - value) & (element > val) then
6:       current - value  $\leftarrow$  element
7:     end if
8:   end for
9: end for
10: return current-value

```

Answer 3.14: The algorithm for finding the Minimum is a brute force over all the elements of the hash table (one list at a time) until we find smallest number in the DS. Because we go over the entire data-structure one item at a time the complexity is $O(n)$

Require: val

```
1: current - value  $\leftarrow \infty$ 
2: for map[0] to map[map.size() - 1] do
3:   list  $\leftarrow$  map[i]
4:   for element in list do
5:     if (element < current - value) then
6:       current - value  $\leftarrow$  element
7:     end if
8:   end for
9: end for
10: return current-value
```

Answer 3.15: The algorithm for finding the Minimum is a brute force over all the elements of the hash table (one list at a time) and count how many numbers are less than or equal val.

Because we go over the entire data-structure one item at a time the complexity is $O(n)$

Require: val

```
1: counter  $\leftarrow$  0
2: for map[0] to map[map.size() - 1] do
3:   list  $\leftarrow$  map[i]
4:   for element in list do
5:     if (element <= val) then
6:       counter  $\leftarrow$  counter + 1
7:     end if
8:   end for
9: end for
10: return current-value
```

Answer 3.16: The algorithm for finding the *i*th number in the hash table is finding the maximum of the DS in $O(n)$ and using the radix sort algorithm from class and going to the *i*th element. Because we go over the entire data-structure to find the maximum in $O(n)$ and then perform radix sort in $O(n)$ the overall time complexity is $O(n)$

4 Designing a data structure according to given specifications

Answer 4.1: sorry don't have enough time :(

"The time is short and the work is a lot, the students are lazy and the Landlord knocks" - pirkey avot

Good Luck!